

Text Mining (Natural Language Processing)

JSC 370: Data Science II

Natural Language Processing (NLP) is used for qualitative data that is collected using open ended or free form text from a survey, medical provider notes in an electronic medical record (EMR), or a transcript of research participant interviews (Koleck et al., 2019).

It is also called 'text mining'.

What is NLP used for?

- Looking at frequencies of words and phrases in text.
- Labeling relationships between words such as subject, object, modification.
- Identify entities in free text, labeling them with types such as person,

location, organization.

- Coupled with AI it can predict words (autocomplete).


How can we do NLP?

- We turn text into numbers.
- Then use R and the tidyverse to explore those numbers.

tidytext: Text mining using dplyr, ggplot2, and other tidy tools

Authors: [Julia Silge](#), [David Robinson](#)

License: [MIT](#)



Using [tidy data principles](#) can make many text mining tasks easier, more effective, and consistent with tools already in wide use. Much of the infrastructure needed for text mining with tidy data frames already exists in packages like [dplyr](#), [broom](#), [tidyr](#) and [ggplot2](#). In this package, we provide functions and supporting data sets to allow conversion of text to and from tidy formats, and to switch seamlessly between tidy tools and existing text mining packages. Check out [our book](#) to learn more about text mining using tidy data principles.

Why tidytext?

Works seamlessly with ggplot2, dplyr and tidyr.

Alternatives:

R: quanteda, tm, koRpus

Python: nltk, Spacy, gensim

Alice's Adventures in Wonderland

Download the alice dataset from [here](#). There are 12 chapters

```
library(tidyverse)
alice <- readRDS("alice.rds")
alice
```

```
## # A tibble: 3,351 × 3
```

```
##   text
```

```
##   <chr>
```

```
## 1 "CHAPTER I."
```

```
## 2 "Down the Rabbit-Hole"
```

```
## 3 ""
```

```
## 4 ""
```

```
## 5 "Alice was beginning to get very tired of sitting by he...
```

```
## 6 "bank, and of having nothing to do: once or twice she h...
```

```
## 7 "the book her sister was reading, but it had no picture...
```

```
## 8 "conversations in it, "and what is the use of a book," ...
```

```
## 9 ""without pictures or conversations?"
```

```
## 10 ""
```

```
chapter chapter_name
```

```
<int> <chr>
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

```
1 CHAPTER I.
```

Tokenizing

Turning text into smaller units, essentially splitting a sentence, phrase, paragraph or entire document into smaller units called tokens (i.e. individual words, numbers, or punctuation marks). Tokenization is needed for natural language processing.

In English:

- split by spaces
- more advanced algorithms

Spacy tokenizer

1. Iterate over whitespace-separated substrings.
2. Look for a token match. If there is a match, stop processing and keep this token.
3. Check whether we have an explicitly defined special case for this substring. If we do, use it.
4. Otherwise, try to consume one prefix. If we consumed a prefix, go back to #2, so that the token match and special cases always get priority.
5. If we didn't consume a prefix, try to consume a suffix and then go back to #2.
6. If we can't consume a prefix or a suffix, look for a URL match.
7. If there's no URL match, then look for a special case.
8. Look for "infixes" — stuff like hyphens etc. and split the substring into tokens on all infixes.
9. Once we can't consume any more of the string, handle it as a single token.

Tokenizing with unnest_tokens

```
library(tidytext)
alice %>%
  unnest_tokens(token, text)

## # A tibble: 26,687 × 3
##   chapter chapter_name token
##   <int>   <chr>         <chr>
## 1       1 1 CHAPTER I.  chapter
## 2       1 1 CHAPTER I.    i
## 3       1 1 CHAPTER I.  down
## 4       1 1 CHAPTER I.  the
## 5       1 1 CHAPTER I.  rabbit
## 6       1 1 CHAPTER I.  hole
## 7       1 1 CHAPTER I.  alice
## 8       1 1 CHAPTER I.  was
## 9       1 1 CHAPTER I.  beginning
## 10      1 1 CHAPTER I.  to
## # ... with 26,677 more rows
```

Words as a unit

Now that we have words as the observation unit we can use the dplyr toolbox.

Using dplyr verbs

```
library(dplyr)
alice %>%
  unnest_tokens(token, text)
```

```
## # A tibble: 26,687 × 3
##   chapter chapter_name token
##   <int> <chr> <chr>
## 1     1 1 CHAPTER I. chapter
## 2     1 1 CHAPTER I. i
## 3     1 1 CHAPTER I. down
## 4     1 1 CHAPTER I. the
## 5     1 1 CHAPTER I. rabbit
## 6     1 1 CHAPTER I. hole
## 7     1 1 CHAPTER I. alice
## 8     1 1 CHAPTER I. was
## 9     1 1 CHAPTER I. beginning
## 10    1 1 CHAPTER I. to
## # ... with 26,677 more rows
```

Using dplyr verbs

```
library(dplyr)
alice %>%
  unnest_tokens(token, text) %>%
  count(token)
```

```
## # A tibble: 2,740 × 2
##   token      n
##   <chr>  <int>
## 1 _alice's    1
## 2 _all        1
## 3 _all_        1
## 4 _and        1
## 5 _are_        4
## 6 _at         1
## 7 _before     1
## 8 _beg_        1
## 9 _began_      1
## 10 _best_      2
## # ... with 2,730 more rows
```

Using dplyr verbs

```
library(dplyr)
alice %>%
  unnest_tokens(token, text) %>%
  count(token, sort = TRUE)
```

```
## # A tibble: 2,740 × 2
##   token      n
##   <chr> <int>
## 1 the    1643
## 2 and     871
## 3 to     729
## 4 a      632
## 5 she    538
## 6 it     527
## 7 of     514
## 8 said   460
## 9 i      393
## 10 alice 386
## # ... with 2,730 more rows
```

Using dplyr verbs

```
library(dplyr)
alice %>%
  unnest_tokens(token, text) %>%
  count(chapter, token)
```

```
## # A tibble: 7,549 × 3
##   chapter token          n
##   <int> <chr>        <int>
## 1     1 1 _curtseying_      1
## 2     2 1 _never_           1
## 3     3 1 _not_             1
## 4     4 1 _one_             1
## 5     5 1 _poison_          1
## 6     6 1 _that_            1
## 7     7 1 _through_         1
## 8     8 1 _took             1
## 9     9 1 _very_            4
## 10    10 1 _was_             1
## # ... with 7,539 more rows
```


Using dplyr verbs

```
library(dplyr)
alice %>%
  unnest_tokens(token, text) %>%
  group_by(chapter) %>%
  count(token) %>%
  top_n(10, n)
```

```
## # A tibble: 122 × 3
## # Groups:   chapter [12]
##   chapter token      n
##   <int> <chr> <int>
## 1     1 1 a      52
## 2     1 1 alice   27
## 3     1 1 and     65
## 4     1 1 i       30
## 5     1 1 it      62
## 6     1 1 of      43
## 7     1 1 she     79
## 8     1 1 the     92
## 9     1 1 to      75
```



Using dplyr verbs and ggplot2

```
library(dplyr)
library(ggplot2)
alice %>%
  unnest_tokens(token, text) %>%
  count(token) %>%
  top_n(10, n) %>%
  ggplot(aes(n, token)) +
  geom_col()
```



Using dplyr verbs and ggplot2

```
library(dplyr)
library(ggplot2)
library(forcats)
alice %>%
  unnest_tokens(token, text) %>%
  count(token) %>%
  top_n(10, n) %>%
  ggplot(aes(n, fct_reorder(token, n))) +
  geom_col()
```



Stop words

A lot of the words don't tell us very much. Words such as "the", "and", "at" and "for" appear a lot in English text but doesn't add much to the context.

Words such as these are called stop words

For more information about differences in stop words and when to remove them read this chapter <https://sm1tar.com/stopwords>

Stop words in tidytext

tidytext comes with a data.frame of stop words

```
stop_words
```

```
## # A tibble: 1,149 × 2
##   word      lexicon
##   <chr>    <chr>
## 1 a       SMART
## 2 a's     SMART
## 3 able    SMART
## 4 about   SMART
## 5 above   SMART
## 6 according SMART
## 7 accordingly SMART
## 8 across   SMART
## 9 actually SMART
## 10 after    SMART
## # ... with 1,139 more rows
```

snowball stopwords

##	[1]	"i"	"me"	"my"	"myself"	"we"
##	[6]	"our"	"ours"	"ourselves"	"you"	"your"
##	[11]	"yours"	"yourself"	"yourselves"	"he"	"him"
##	[16]	"his"	"himself"	"she"	"her"	"hers"
##	[21]	"herself"	"it"	"its"	"itself"	"they"
##	[26]	"them"	"their"	"theirs"	"themselves"	"what"
##	[31]	"which"	"who"	"whom"	"this"	"that"
##	[36]	"these"	"those"	"am"	"is"	"are"
##	[41]	"was"	"were"	"be"	"been"	"being"
##	[46]	"have"	"has"	"had"	"having"	"do"
##	[51]	"does"	"did"	"doing"	"would"	"should"
##	[56]	"could"	"ought"	"i'm"	"you're"	"he's"
##	[61]	"she's"	"it's"	"we're"	"they're"	"i've"
##	[66]	"you've"	"we've"	"they've"	"i'd"	"you'd"
##	[71]	"he'd"	"she'd"	"we'd"	"they'd"	"i'll"
##	[76]	"you'll"	"he'll"	"she'll"	"we'll"	"they'll"
##	[81]	"isn't"	"aren't"	"wasn't"	"weren't"	"hasn't"
##	[86]	"haven't"	"hadn't"	"doesn't"	"don't"	"didn't"
##	[91]	"won't"	"wouldn't"	"shan't"	"shouldn't"	"can't"
##	[96]	"cannot"	"couldn't"	"mustn't"	"let's"	"that's"

Removing stopwords

We can use an `anti_join()` to remove the tokens that also appear in the `stop_words` data.frame

```
alice %>%  
  unnest_tokens(token, text) %>%  
  anti_join(stop_words, by = c("token" = "word")) %>%  
  count(token, sort = TRUE)
```

```
## # A tibble: 2,314 × 2  
##   token      n  
##   <chr>  <int>  
## 1 alice   386  
## 2 time    71  
## 3 queen   68  
## 4 king    61  
## 5 don't   60  
## 6 it's    57  
## 7 i'm     56  
## 8 mock    56
```


Anti-join with same variable name

```
alice %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words, by = c("word")) %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 2,314 × 2  
##   word      n  
##   <chr>  <int>  
## 1 alice   386  
## 2 time    71  
## 3 queen   68  
## 4 king    61  
## 5 don't   60  
## 6 it's    57  
## 7 i'm     56  
## 8 mock    56  
## 9 turtle  56  
## 10 gryphon 55  
## # ... with 2,304 more rows
```

Stop words removed

```
alice %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words, by = c("word")) %>%  
  count(word, sort = TRUE) %>%  
  top_n(10, n) %>%  
  ggplot(aes(n, fct_reorder(word, n))) +  
  geom_col()
```



Wordcloud

```
library(wordcloud)
pal<-brewer.pal(8,"Spectral")
alice %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words, by = c("word")) %>%
  count(word, sort = TRUE) %>%
  top_n(10, n) %>%
  with(wordcloud(word, n, random.order = FALSE, max.words = 100, colors=pal))
```



Which words appear together?

ngrams are n consecutive word, we can count these to see what words appears together.

- ngram with $n = 1$ are called unigrams: "which", "words", "appears", "together"
- ngram with $n = 2$ are called bigrams: "which words", "words appears", "appears together"
- ngram with $n = 3$ are called trigrams: "which words appears", "words appears together"

Which words appears together?

We can extract bigrams using `unnest_ngrams()` with `n = 2`

```
alice %>%
  unnest_ngrams(ngram, text, n = 2)

## # A tibble: 25,170 × 3
##   chapter chapter_name ngram
##   <int> <chr>         <chr>
## 1     1 1 CHAPTER I.    chapter i
## 2     1 1 CHAPTER I.    down the
## 3     1 1 CHAPTER I.    the rabbit
## 4     1 1 CHAPTER I.    rabbit hole
## 5     1 1 CHAPTER I.    <NA>
## 6     1 1 CHAPTER I.    <NA>
## 7     1 1 CHAPTER I.    alice was
## 8     1 1 CHAPTER I.    was beginning
## 9     1 1 CHAPTER I.    beginning to
## 10    1 1 CHAPTER I.    to get
## # ... with 25,160 more rows
```

Which words appears together?

Tallying up the bi-grams still shows a lot of stop words but is able to pick up relationships with patients

```
alice %>%  
  unnest_ngrams(ngram, text, n = 2) %>%  
  count(ngram, sort = TRUE)
```

```
## # A tibble: 13,424 x 2  
##   ngram      n  
##   <chr>    <int>  
## 1 <NA>      951  
## 2 said the  206  
## 3 of the   130  
## 4 said alice 112  
## 5 in a     96  
## 6 and the   75  
## 7 in the    75  
## 8 it was    72
```

Which words appears together?

```
alice %>%
  unnest_ngrams(ngram, text, n = 2) %>%
  separate(ngram, into = c("word1", "word2"), sep = " ") %>%
  select(word1, word2)
```

```
## # A tibble: 25,170 × 2
##   word1      word2
##   <chr>    <chr>
## 1 chapter  i
## 2 down    the
## 3 the     rabbit
## 4 rabbit  hole
## 5 <NA>    <NA>
## 6 <NA>    <NA>
## 7 alice   was
## 8 was     beginning
## 9 beginning to
## 10 to     get
## # ... with 25,160 more rows
```

```
alice %>%
  unnest_ngrams(ngram, text, n = 2) %>%
  separate(ngram, into = c("word1", "word2"), sep = " ") %>%
  select(word1, word2) %>%
  filter(word1 == "alice")
```

```
## # A tibble: 336 × 2
##   word1 word2
##   <chr> <chr>
## 1 alice was
## 2 alice think
## 3 alice started
## 4 alice after
## 5 alice had
## 6 alice to
## 7 alice had
## 8 alice had
## 9 alice soon
## 10 alice began
## # ... with 326 more rows
```

```
alice %>%
  unnest_ngrams(ngram, text, n = 2) %>%
  separate(ngram, into = c("word1", "word2"), sep = " ") %>%
  select(word1, word2) %>%
  filter(word1 == "alice") %>%
  count(word2, sort = TRUE)
```

```
## # A tibble: 133 × 2
##   word2      n
##   <chr>  <int>
## 1 and      18
## 2 was      17
## 3 thought  12
## 4 as       11
## 5 said     11
## 6 could    10
## 7 had      10
## 8 did       9
## 9 in        9
## 10 to       9
## # ... with 123 more rows
```

```
alice %>%
  unnest_ngrams(ngram, text, n = 2) %>%
  separate(ngram, into = c("word1", "word2"), sep = " ") %>%
  select(word1, word2) %>%
  filter(word2 == "alice") %>%
  count(word1, sort = TRUE)
```

```
## # A tibble: 106 × 2
##   word1      n
##   <chr>   <int>
## 1 said    112
## 2 thought  25
## 3 to      22
## 4 and     15
## 5 poor    11
## 6 cried    7
## 7 at       6
## 8 so       6
## 9 that     5
## 10 exclaimed 3
## # ... with 96 more rows
```


TF-IDF

TF: Term frequency gives weight to terms that appear a lot. It's a measure of how important a word may be and how frequently a word occurs within a document (e.g. a book chapter). IDF decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents (e.g. all chapters in a book).

Some words that occur many times in a document may not be important; in English, these are probably words like “the”, “is”, “of”, and so forth. We might take the approach of adding words like these to a list of stop words and removing them before analysis, but it is possible that some of these words might be more important in some documents than others. A list of stop words is not a sophisticated approach to adjusting term frequency for commonly used words.

TF-IDF

IDF: Inverse document frequency

IDF decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents.

The inverse document frequency for any given term is defined as

$$idf(term) = \ln\left(\frac{\text{n documents}}{\text{n documents containing term}}\right)$$

TF-IDF

TF-IDF: TF and IDF can be combined (the two quantities multiplied together), which is the frequency of a term adjusted for how rarely it is used.

The idea of TF-IDF is to find the important words for the content of each document by decreasing the weight for commonly used words and increasing the weight for words that are not used very much in a collection or corpus of documents.

TF-IDF with tidytext

```
alice %>%  
  unnest_tokens(text, text)
```

```
## # A tibble: 26,687 × 3  
##   text      chapter chapter_name  
##   <chr>      <int> <chr>  
## 1 chapter          1 CHAPTER I.  
## 2 i                1 CHAPTER I.  
## 3 down            1 CHAPTER I.  
## 4 the             1 CHAPTER I.  
## 5 rabbit          1 CHAPTER I.  
## 6 hole            1 CHAPTER I.  
## 7 alice           1 CHAPTER I.  
## 8 was             1 CHAPTER I.  
## 9 beginning       1 CHAPTER I.  
## 10 to             1 CHAPTER I.  
## # ... with 26,677 more rows
```

TF-IDF with tidytext

```
alice %>%  
  unnest_tokens(text, text) %>%  
  count(text, chapter)
```

```
## # A tibble: 7,549 × 3  
##   text      chapter     n  
##   <chr>      <int> <int>  
## 1 _alice's      2     1  
## 2 _all          12     1  
## 3 _all_         12     1  
## 4 _and          9     1  
## 5 _are_         4     1  
## 6 _are_         6     1  
## 7 _are_         8     1  
## 8 _are_         9     1  
## 9 _at           9     1  
## 10 _before      12     1  
## # ... with 7,539 more rows
```

TF-IDF with tidytext

```
alice %>%  
  unnest_tokens(text, text) %>%  
  count(text, chapter) %>%  
  bind_tf_idf(text, chapter, n)
```

```
## # A tibble: 7,549 × 6  
##   text      chapter      n      tf      idf      tf_idf  
##   <chr>      <int> <int>    <dbl> <dbl>    <dbl>  
## 1 _alice's      2      1 0.000471 2.48 0.00117  
## 2 _all          12      1 0.000468 2.48 0.00116  
## 3 _all_         12      1 0.000468 2.48 0.00116  
## 4 _and          9      1 0.000435 2.48 0.00108  
## 5 _are_         4      1 0.000375 1.10 0.000411  
## 6 _are_         6      1 0.000382 1.10 0.000420  
## 7 _are_         8      1 0.000400 1.10 0.000439  
## 8 _are_         9      1 0.000435 1.10 0.000478  
## 9 _at           9      1 0.000435 2.48 0.00108  
## 10 _before     12      1 0.000468 2.48 0.00116  
## # ... with 7,539 more rows
```

TF-IDF with tidytext

```
alice %>%
  unnest_tokens(text, text) %>%
  count(text, chapter) %>%
  bind_tf_idf(text, chapter, n) %>%
  arrange(desc(tf_idf))
```

```
## # A tibble: 7,549 × 6
##   text      chapter     n      tf      idf tf_idf
##   <chr>      <int> <int>   <dbl> <dbl>   <dbl>
## 1 dormouse      7     26 0.0112  1.79 0.0201
## 2 hatter        7     32 0.0138  1.39 0.0191
## 3 mock         10     28 0.0136  1.39 0.0189
## 4 turtle        10     28 0.0136  1.39 0.0189
## 5 gryphon       10     31 0.0151  1.10 0.0166
## 6 turtle         9     27 0.0117  1.39 0.0163
## 7 caterpillar   5     25 0.0115  1.39 0.0159
## 8 dance        10     13 0.00632 2.48 0.0157
## 9 mock          9     26 0.0113  1.39 0.0157
## 10 hatter       11     21 0.0110  1.39 0.0153
```

Sentiment Analysis

- Sentiment Analysis is a process of extracting opinions that have different scores like positive, negative or neutral.
- Based on sentiment analysis, you can find out the nature of opinion or sentences in text.
- Sentiment Analysis is a type of classification where the data are classified into different classes like positive or negative or happy, sad, angry, etc.

Sentiment Analysis

```
positive <- get_sentiments("bing") %>%
  filter(sentiment == "positive")

alice %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words, by = c("word")) %>%
  semi_join(positive) %>%
  count(word, sort = TRUE)

## # A tibble: 140 × 2
##   word      n
##   <chr>  <int>
## 1 beautiful 13
## 2 majesty   12
## 3 glad      11
## 4 bright     8
## 5 eagerly    8
## 6 ready       8
## 7 top         8
```

Sentiment Analysis

```
bing <- get_sentiments("bing")
alicesentiment<-alice %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words, by = c("word")) %>%
  inner_join(bing) %>%
  count(word, sentiment, sort = TRUE)
alicesentiment
```

```
## # A tibble: 413 × 3
##   word      sentiment      n
##   <chr>    <chr>    <int>
## 1 mock      negative     56
## 2 poor      negative     27
## 3 hastily   negative     16
## 4 mad        negative     15
## 5 anxiously negative     14
## 6 beautiful positive     13
## 7 afraid     negative     12
## 8 majesty    positive     12
```

Sentiment Analysis

```
alicesentiment %>%  
  filter(n > 7) %>%  
  mutate(n = ifelse(sentiment == "negative", -n, n)) %>%  
  mutate(word = reorder(word, n)) %>%  
  ggplot(aes(word, n, fill = sentiment)) +  
  geom_col() +  
  coord_flip() +  
  labs(y = "Contribution to sentiment")
```

